

Signed p -adic Residual Encodings of Finite-Domain All-Different Systems

with a Sudoku Case Study*

Greg Baker

School of Computing, Australian National University

Canberra, Australia

greg.baker@anu.edu.au

ORCID: 0000-0002-6910-3011

Abstract

We encode finite-domain all-different systems as signed synthetic p -adic residual objectives, with Sudoku as a worked example. Given a graph whose vertices carry finite allowed sets and whose edges require unequal assignments, positive synthetic observations create unary domain wells and negative synthetic observations reward pairwise inequality. For primes $p > q$, every nonzero difference between values in $\{1, \dots, q\}$ has p -adic norm 1. On domain-respecting states, the objective is therefore an additive constant plus the edge-conflict count. Its global minimisers are exactly the domain-respecting assignments that maximise the number of unequal edges. When the associated list-colouring instance is satisfiable, these minimisers are precisely its proper list-colourings.

Standard Sudoku is a direct instance: there is one variable per cell, singleton domains on clue cells, and edges between cells that share a row, column, or box. The construction uses the original 81 cell variables rather than a one-hot lift to 729 Boolean indicators. The Sudoku section is not a solver proposal; it spells out the encoding in a case where the variables, domains, and all-different constraints are familiar. The word “regression” refers to the affine residual form of the objective, not to an ordinary positive-weight statistical loss.

Keywords: p -adic regression; signed residual objectives; all-different constraints; list colouring; Sudoku

1 Introduction and prior work

This paper encodes finite-domain all-different systems as signed synthetic p -adic residual objectives, with Sudoku as a worked example. The representation keeps one variable per semantic object, uses the same finite domain as the original problem, and enforces both domain membership and pairwise distinctness in one objective.

All signed objectives in this paper are optimised over $\mathbb{Z}_p^n$. On this domain every affine residual is bounded by 1. Once $p > q$, the same residual also becomes an exact 0/1 equality indicator on the intended finite alphabet. The restriction to $\mathbb{Z}_p^n$ is mainly a modelling choice, not a way to hide an unbounded objective: although a single negative term $-|x_i - x_j|_p$ is unbounded below over $\mathbb{Q}_p^n$, the objectives constructed here remain bounded below on $\mathbb{Q}_p^n$. The same domination condition that forces digit-snapping, $\lambda > \Delta$, also dominates the total inequality reward.

Two earlier results motivate the construction:

*This article presents material that also appears as a chapter of the author’s PhD dissertation (Australian National University, in preparation); it is the standalone version of that work.

- A structural result on p -adic linear regression [1]: an optimal hyperplane is supported by $n + 1$ data points.
- A complexity result on p -adic linear regression [2]: the problem is NP-hard, via an embedding of MAX-CUT using positive pinning points that snap coefficients to $\{0, 1\}$.

Recent work in *p-Adic Numbers, Ultrametric Analysis and Applications* connects p -adic methods with regression structure and optimisation [1, 7]. Here signed residual terms enforce pairwise distinctness across finite domains. The main theorem turns a finite-domain all-different or list-colouring instance into a signed weighted synthetic p -adic residual objective whose global minimisers are exactly the domain-respecting assignments with the fewest conflicts, and in the satisfiable case exactly the valid solutions. Relative to [1, 2], this replaces structural support results and binary forcing gadgets by a polynomial-time finite-domain reduction with an exact global-minimiser description.

Negative signed residual terms do not have the usual interpretation of Euclidean least-squares terms. Classical generalized ridge regression allowed negative ridge parameters [9]; later work found zero or negative risk-optimal ridge parameters in overparameterised models [10] and under distribution shift [11]. Related anti-shrinkage effects appear in Gaussian model-uncertainty formulations [12] and in multivariate climate downscaling [13]. These results do not justify the present construction by analogy. They only show that negative quadratic penalties are not, by themselves, incoherent. In the p -adic construction below, the reason they are usable is sharper: on $\mathbb{Z}_p^n$, the residuals are bounded, and on the intended finite domain they become equality indicators.

From the constraint-programming side, the construction is close to weighted CSP / Max-CSP penalty functions, graph-colouring energy models, and Ising/Potts/QUBO encodings of discrete optimisation. Sudoku itself is routinely formulated using exact cover, SAT, or *alldifferent* constraints [3, 4, 5]. The feature specific to this paper is the affine p -adic residual form. Once p is larger than the finite domain and the unary wells snap the variables to their allowed values, each residual $|x_i - x_j|_p$ is an equality/inequality indicator on the original domain. The same mechanism also handles clauses: after Boolean pinning, a 3-CNF clause is one forbidden affine equality, so a signed residual can reward its satisfaction.

2 Preliminaries

2.1 p -adic integers and norms

Fix a prime p . The p -adic valuation $v_p(n)$ of an integer $n \neq 0$ is the exponent of p in its factorisation, and the p -adic absolute value is

$$|n|_p = p^{-v_p(n)}, \quad |0|_p = 0.$$

We will work in the p -adic integers $\mathbb{Z}_p$, where all norms satisfy $|x|_p \leq 1$. The key choice for Sudoku is to take $p > 9$ (e.g. $p = 11$), so that the p -adic absolute value of any Sudoku digit or digit difference is always 0 or 1.

For comparison, note that the discrete metric $d_{\text{disc}}(s, t) := \mathbf{1}_{s \neq t}$ is also an ultrametric. We return to a companion discrete-metric version of the reduction after the main correctness theorem.

2.2 List colouring

Many finite-domain constraint systems built from all-different conditions can be viewed as list-colouring problems on a graph. Each variable x_i is assigned a value from a finite allowed set D_i , and an edge $\{i, j\}$ expresses the preference or requirement that $x_i \neq x_j$. Thus an all-different

constraint of arity r contributes the $\binom{r}{2}$ pairwise edges of its primal-graph clique; this paper uses that pairwise primal graph, not global *alldifferent* propagation.

Figure 1 gives a concrete toy instance.

One solution is $x = (1, 2, 3)$.

2.3 Residual-objective formulation

The equivalent dataset for the toy instance in Figure 1 is shown in Figure 2.

Its synthetic dataset has six positive observations, one for each allowed value in the three domains, and two negative observations, one for each edge. At the coefficient vector $x = (1, 2, 3)$, every edge reward is active and the unary term attains its baseline value.

3 A general all-different residual-objective template

Over $\mathbb{Z}_p$ with $p > q$, once values are restricted to $\{1, \dots, q\}$, every nonzero difference has norm 1.

Definition 1 (Signed synthetic p -adic residual dataset). *For parameters $x \in \mathbb{Z}_p^n$, a signed weighted affine observation is a quadruple (σ, u, b, γ) , where $\sigma \in \{+1, -1\}$ is the sign, $u \in \mathbb{Z}^n$ is the feature vector, $b \in \mathbb{Z}$ is the target, and $\gamma > 0$ is the weight. It contributes the term*

$$\sigma \gamma |u^\top x - b|_p$$

to the loss; formally, this is the p -adic analogue of a weighted Euclidean residual term.

Throughout this paper x ranges over $\mathbb{Z}_p^n$. Since $u \in \mathbb{Z}^n$, $b \in \mathbb{Z}$, and $x \in \mathbb{Z}_p^n$, each residual $u^\top x - b$ lies in $\mathbb{Z}_p$, and hence $|u^\top x - b|_p \leq 1$. Every finite signed objective considered here is therefore bounded on its stated domain.

We will use the phrase *signed p -adic residual objective* for sums of such terms. The word “regression” refers to the affine residual form of the objective; because negative weights are allowed, the resulting functional is an optimisation objective rather than a statistical loss in the usual positive-weight sense.

Definition 2 (Positive/negative observations). *A signed weighted affine observation (σ, u, b, γ) is positive if $\sigma = +1$ and negative if $\sigma = -1$. In the constructions below, positive observations serve as pinning terms and negative observations serve as reward terms.*

Theorem 1 (Finite-domain signed affine compiler template). *Fix a prime p . Let $D_i \subseteq \mathbb{Z}$ be finite nonempty sets whose elements are pairwise incongruent modulo p , and put*

$$D := \prod_{i=1}^n D_i \subseteq \mathbb{Z}_p^n.$$




Figure 1: A tiny list-colouring instance. The assignment $x_1 = 1$, $x_2 = 2$, $x_3 = 3$ respects the domains and makes every edge unequal.

Type	Synthetic observations	Role
positive	$(+1, e_1, 1, 3)$	pin x_1 to D_1
positive	$(+1, e_1, 2, 3)$	pin x_1 to D_1
positive	$(+1, e_2, 2, 3)$	pin x_2 to D_2
positive	$(+1, e_2, 3, 3)$	pin x_2 to D_2
positive	$(+1, e_3, 1, 3)$	pin x_3 to D_3
positive	$(+1, e_3, 3, 3)$	pin x_3 to D_3
negative	$(-1, e_1 - e_2, 0, 1)$	reward $x_1 \neq x_2$
negative	$(-1, e_2 - e_3, 0, 1)$	reward $x_2 \neq x_3$

Figure 2: Synthetic dataset for the toy list-colouring instance in Figure 1. Here e_1, e_2, e_3 denote the standard basis vectors of $\mathbb{Z}^3$.

For each coordinate i , choose a pinning weight $\lambda_i > 0$. Let $\mathcal{T}$ be a finite family of signed affine observations

$$(\sigma_\ell, u_\ell, b_\ell, \gamma_\ell), \quad \sigma_\ell \in \{\pm 1\}, \quad u_\ell \in \mathbb{Z}^n, \quad b_\ell \in \mathbb{Z}, \quad \gamma_\ell > 0,$$

and define

$$L(x) := \sum_{i=1}^n \lambda_i \sum_{a \in D_i} |x_i - a|_p + \sum_{\ell \in \mathcal{T}} \sigma_\ell \gamma_\ell |u_\ell^\top x - b_\ell|_p.$$

Assume that, for every coordinate i ,

$$\lambda_i > \sum_{\ell: (u_\ell)_i \neq 0} \gamma_\ell,$$

and that, for every $x \in D$ and every $\ell \in \mathcal{T}$, the residual $u_\ell^\top x - b_\ell$ is either 0 or a p -adic unit. Then L attains its global minimum on D , and every global minimiser lies in D . On D , L is a constant plus a finite signed objective whose interaction terms have values 0 and $\sigma_\ell \gamma_\ell$.

Proof. Suppose $x_i \notin D_i$. If x_i is congruent modulo p to an element of D_i , let b be that element; otherwise choose any $b \in D_i$. Let x' be obtained from x by replacing x_i with b .

The i -th unary sum decreases by at least $\lambda_i |x_i - b|_p$. Indeed, if $x_i \equiv b \pmod{p}$, then the only changed unary contribution below 1 is $|x_i - b|_p$. If x_i is congruent to no element of D_i , then $|x_i - a|_p = 1$ for all $a \in D_i$, while the unary sum at b has one zero term and all other terms equal to 1; in this case $|x_i - b|_p = 1$.

For an interaction term with $(u_\ell)_i \neq 0$, write $r'_\ell = u_\ell^\top x' - b_\ell$. Then

$$u_\ell^\top x - b_\ell = r'_\ell + (u_\ell)_i (x_i - b).$$

Since $(u_\ell)_i \in \mathbb{Z}$, $|(u_\ell)_i|_p \leq 1$, and the ultrametric inequality gives

$$||u_\ell^\top x - b_\ell|_p - |r'_\ell|_p| \leq |x_i - b|_p.$$

Thus the total interaction increase caused by changing coordinate i is at most

$$\sum_{\ell: (u_\ell)_i \neq 0} \gamma_\ell |x_i - b|_p.$$

By the assumed domination inequality, $L(x') < L(x)$. Repeating this coordinate-snapping step gives a point of D with smaller loss than any starting point outside D . Since D is finite, the global minimum is attained on D , and no global minimiser lies outside D .

On D , every unary sum is constant: exactly one term is zero and the rest are units. The second assumption says each interaction residual is either zero or a unit, so each signed interaction term is either 0 or $\sigma_\ell \gamma_\ell$. This proves the stated finite-domain form of $L|_D$. $\square$

Definition 3 (Synthetic dataset for list-colouring). Let $V = \{1, \dots, n\}$, let $G = (V, E)$ be a finite graph, let $D_i \subseteq \{1, \dots, q\}$ be nonempty allowed sets, and fix a prime $p > q$. Define

$$w_{ii} := 0, \quad w_{ij} := \begin{cases} 1, & \{i, j\} \in E, \\ 0, & \{i, j\} \notin E, \end{cases} \quad (i \neq j),$$

let $\Delta := \max_{i \in V} \sum_{j=1}^n w_{ij}$, and set $\lambda := 1 + \Delta$. Writing $e_1, \dots, e_n$ for the standard basis vectors of $\mathbb{Z}^n$, define

$$\mathcal{S}^+ := \{(+1, e_i, a, \lambda) : i \in V, a \in D_i\},$$

$$\mathcal{S}^- := \{(-1, e_i - e_j, 0, 1) : 1 \leq i < j \leq n, w_{ij} = 1\}.$$

Set $\mathcal{S} := \mathcal{S}^+ \sqcup \mathcal{S}^-$.

The theorem is stated for a simple graph, so $w_{ij} \in \{0, 1\}$. An edge-weighted version is obtained by replacing the unit negative observation on edge $\{i, j\}$ with weight $c_{ij} > 0$ and choosing the pinning weight larger than $\max_i \sum_j c_{ij}$, but the binary case is the one needed here.

Definition 4 (Loss associated with the list-colouring dataset). For $\mathcal{S}$ as in Definition 3, define

$$L_{\mathcal{S}}(x) = \lambda \sum_{i=1}^n \sum_{a \in D_i} |x_i - a|_p - \sum_{1 \leq i < j \leq n} w_{ij} |x_i - x_j|_p.$$

This dataset has exactly $\sum_{i=1}^n |D_i|$ positive observations and $|E|$ negative observations, so it is computable in time $O(\sum_i |D_i| + |E|)$. The positive observations pin coordinates to allowed values, and the negative observations reward inequality across edges.

Each positive observation $(+1, e_i, a, \lambda)$ contributes $\lambda |x_i - a|_p$ to the loss, and each negative observation $(-1, e_i - e_j, 0, 1)$ contributes $-|x_i - x_j|_p$ to the loss. Summing them yields the displayed loss. The size bound and runtime are immediate from the explicit definitions.

The compiler template already implies that minimisers snap to the finite alphabet, because the pinning weight is larger than the degree of every vertex. The next proof records the stronger fact needed for all-different systems: after snapping, the objective is exactly a constant plus the edge-conflict count.

3.1 Correctness

Since $p > q$, the residues of $1, \dots, q$ are distinct modulo p . In particular, for any fixed $i \in V$ and any $t \in \mathbb{Z}_p$, there is at most one $a \in D_i$ with $t \equiv a \pmod{p}$.

First we record the simple observation that, on $\{1, \dots, q\}$, the p -adic norm detects inequality exactly.

Lemma 1 (Edge terms become inequality indicators on the domain). If $x_i, x_j \in \{1, \dots, q\}$ and $p > q$, then

$$|x_i - x_j|_p = \begin{cases} 0, & x_i = x_j, \\ 1, & x_i \neq x_j. \end{cases}$$

Proof. If $x_i = x_j$ then the difference is zero. If $x_i \neq x_j$, then $1 \leq |x_i - x_j| \leq q - 1 < p$, so $p \nmid (x_i - x_j)$ and therefore $|x_i - x_j|_p = 1$. $\square$

We next prove a digit-snapping lemma: if x minimises $L_{\mathcal{S}}$, then each coordinate x_i lies in D_i . The explicit unary-well formula makes that argument transparent.

Lemma 2 (Unary wells). *For $i \in V$, define*

$$U_i(t) := \sum_{a \in D_i} |t - a|_p \quad (t \in \mathbb{Z}_p).$$

Then $U_i(t) \geq |D_i| - 1$, with equality if and only if $t \in D_i$.

Proof. If $t \in D_i$, then one summand in $U_i(t)$ is 0, and the remaining $|D_i| - 1$ summands are 1, so $U_i(t) = |D_i| - 1$.

If $t \notin D_i$ but $t \equiv b \pmod{p}$ for some $b \in D_i$, then $0 < |t - b|_p < 1$, while for every $a \in D_i \setminus \{b\}$ we have $t - a \equiv b - a \not\equiv 0 \pmod{p}$, so $|t - a|_p = 1$. Hence

$$U_i(t) = (|D_i| - 1) + |t - b|_p > |D_i| - 1.$$

If t is congruent modulo p to no element of D_i , then $|t - a|_p = 1$ for all $a \in D_i$, so $U_i(t) = |D_i| > |D_i| - 1$. $\square$

Lemma 3 (Digit-snapping lemma). *Let $x \in \mathbb{Z}_p^n$. Suppose there exists $i \in V$ with $x_i \notin D_i$. Choose $b \in D_i$ as follows: if x_i is congruent modulo p to an element of D_i , take that unique element; otherwise choose any $b \in D_i$. Let x' be obtained from x by replacing x_i with b . Then*

$$L_S(x') < L_S(x).$$

Proof. The positive part of the loss changes by

$$\lambda \left(\sum_{a \in D_i} |b - a|_p - \sum_{a \in D_i} |x_i - a|_p \right).$$

There are two cases to consider.

If x_i is congruent modulo p to no element of D_i , then $|x_i - a|_p = 1$ for every $a \in D_i$, so $\sum_{a \in D_i} |x_i - a|_p = |D_i|$. On the other hand, one summand in $\sum_{a \in D_i} |b - a|_p$ is 0 and the remaining $|D_i| - 1$ summands are 1, so $\sum_{a \in D_i} |b - a|_p = |D_i| - 1$. Hence the positive change is $-\lambda$. Since $|x_i - b|_p = 1$ in this case, this is at most $-\lambda|x_i - b|_p$.

If $x_i \equiv b \pmod{p}$ for some $b \in D_i$, then $\sum_{a \in D_i} |b - a|_p = |D_i| - 1$, whereas

$$\sum_{a \in D_i} |x_i - a|_p = (|D_i| - 1) + |x_i - b|_p.$$

Hence the positive change is $-\lambda|x_i - b|_p$.

For each $j \in V$, the corresponding negative component changes by

$$w_{ij}(|x_i - x_j|_p - |b - x_j|_p).$$

By the ultrametric inequality,

$$|x_i - x_j|_p \leq \max(|x_i - b|_p, |b - x_j|_p),$$

so

$$|x_i - x_j|_p - |b - x_j|_p \leq |x_i - b|_p.$$

Summing over all $j \in V$,

$$L_S(x') - L_S(x) \leq \left(-\lambda + \sum_{j=1}^n w_{ij} \right) |x_i - b|_p \leq -|x_i - b|_p < 0.$$

$\square$

On domain-respecting assignments, the loss is an additive constant plus the conflict count.

Theorem 2 (Correctness of the synthetic reduction). *For the dataset $\mathcal{S} = \mathcal{S}^+ \sqcup \mathcal{S}^-$ from Definition 3 and the loss from Definition 4, every global minimiser of $L_{\mathcal{S}}$ lies in $\prod_{i=1}^n D_i$. On $\prod_{i=1}^n D_i$,*

$$L_{\mathcal{S}}(x) = \lambda \sum_{i=1}^n (|D_i| - 1) - |E| + \sum_{1 \leq i < j \leq n} w_{ij} \mathbf{1}_{x_i = x_j}.$$

Consequently, the global minimisers are exactly the domain-respecting assignments that minimise the conflict sum $\sum_{1 \leq i < j \leq n} w_{ij} \mathbf{1}_{x_i = x_j}$, or equivalently maximise the number of unequal edges. In particular, if the list-colouring instance $(G, (D_i)_{i=1}^n)$ is satisfiable, then this minimum conflict sum is 0, and the global minimisers are exactly its proper list-colourings.

Proof. By Lemma 3, any point with some coordinate outside its domain can be improved by snapping that coordinate into the domain. Therefore every global minimiser lies in $\prod_i D_i$. If $x \in \prod_i D_i$, then Lemma 2 gives $U_i(x_i) = |D_i| - 1$ for every i , and Lemma 1 gives

$$|x_i - x_j|_p = \mathbf{1}_{x_i \neq x_j} \quad (1 \leq i < j \leq n).$$

Since $\mathbf{1}_{x_i \neq x_j} = 1 - \mathbf{1}_{x_i = x_j}$ on every edge, substituting these identities into $L_{\mathcal{S}}$ yields the displayed formula. A global minimiser exists: $L_{\mathcal{S}}$ takes only finitely many values on the finite set $\prod_i D_i$, and by the previous paragraph no point outside $\prod_i D_i$ can be a global minimiser, so the global minimum is attained on $\prod_i D_i$. There the displayed formula is an additive constant plus the conflict sum, so a domain-respecting assignment is a global minimiser if and only if it attains the least possible conflict sum; this is the stated “exactly” characterisation in both directions. In particular, if $(G, (D_i)_{i=1}^n)$ is satisfiable the least conflict sum is 0, and the global minimisers are exactly its proper list-colourings. $\square$

Corollary 3. *Any finite-domain all-different constraint system can be encoded, in polynomial time, as a signed weighted synthetic p -adic residual objective whose global minimisers are exactly the domain-respecting assignments with minimum conflict count. If the constraint system is satisfiable, these minimisers are exactly its satisfying assignments.*

Proof. Given an all-different constraint system, form its primal graph: the vertices are the variables, and two variables are adjacent if they appear together in some all-different constraint. The satisfying assignments are exactly the proper list-colourings of this graph, with the variable domains as lists, and the minimum-conflict assignments are exactly the domain-respecting assignments minimising the corresponding edge-conflict sum. Apply Definitions 3 and 4, and Theorem 2. $\square$

Discrete-metric variant. The same synthetic dataset also works with the discrete metric

$$d_{\text{disc}}(s, t) := \mathbf{1}_{s \neq t}.$$

If each residual $|u^\top x - b|_p$ is replaced by $d_{\text{disc}}(u^\top x, b)$, then the loss becomes

$$L_{\mathcal{S}}^{\text{disc}}(x) = \lambda \sum_{i=1}^n \sum_{a \in D_i} \mathbf{1}_{x_i \neq a} - \sum_{1 \leq i < j \leq n} w_{ij} \mathbf{1}_{x_i \neq x_j}.$$

The unary wells satisfy

$$U_i^{\text{disc}}(t) := \sum_{a \in D_i} \mathbf{1}_{t \neq a} = \begin{cases} |D_i| - 1, & t \in D_i, \\ |D_i|, & t \notin D_i. \end{cases}$$

Hence replacing any $x_i \notin D_i$ by a value $b \in D_i$ decreases the positive part by exactly λ , while the total increase in the negative part is at most $\sum_j w_{ij} \leq \Delta$. Since $\lambda = 1 + \Delta$, the same snapping argument goes through. Therefore the conclusions of Theorem 2 and Corollary 3 remain valid for the discrete metric as well. The p -adic case remains the main one because it realises the same indicator behaviour inside the affine p -adic regression framework used throughout the paper.

Boolean clause rewards. The same signed-residual idea also applies beyond all-different constraints. Let

$$\Phi = \bigwedge_{r=1}^m C_r$$

be a 3-CNF formula on variables $z_1, \dots, z_n$, and encode truth values by

$$x_i = 0 \iff z_i \text{ is true}, \quad x_i = 1 \iff z_i \text{ is false}.$$

For a literal λ on variable z_i , define its failure indicator by

$$\delta(\lambda; x) = \begin{cases} x_i, & \lambda = z_i, \\ 1 - x_i, & \lambda = \neg z_i. \end{cases}$$

Then for a clause $C = (\lambda_1 \vee \lambda_2 \vee \lambda_3)$,

$$C(x) \text{ is false} \iff \delta(\lambda_1; x) + \delta(\lambda_2; x) + \delta(\lambda_3; x) = 3.$$

After moving constants to the right-hand side, this means that there are $u_C \in \mathbb{Z}^n$ and $t_C \in \mathbb{Z}$ such that

$$C(x) \text{ is true} \iff u_C^\top x \neq t_C.$$

Here u_C records the ± 1 literal coefficients, and $t_C = 3 - k$, where k is the number of negated literals in the clause. The translation is mechanical: a positive literal contributes $+x_i$, a negated literal contributes $-x_i$, and the forbidden right-hand side is $3 - k$. For example,

$$(z_1 \vee z_2 \vee \neg z_3) \rightsquigarrow x_1 + x_2 - x_3 \neq 2.$$

If $p > 3$ and $x \in \{0, 1\}^n$, then $u_C^\top x - t_C$ is 0 when C is false and one of $-1, -2, -3$ when C is true, so

$$|u_C^\top x - t_C|_p = \mathbf{1}_{C(x) \text{ is true}}.$$

Thus a negative observation $(-1, u_C, t_C, 1)$ rewards satisfaction of the clause. If Δ is the maximum number of clauses containing any variable and $\alpha > \Delta$, then

$$L_\Phi(x) = \alpha \sum_{i=1}^n (|x_i|_p + |x_i - 1|_p) - \sum_{C \in \Phi} |u_C^\top x - t_C|_p$$

has the same snapping property as before: every global minimiser lies in $\{0, 1\}^n$, and on Boolean assignments

$$L_\Phi(x) = \alpha n - \text{sat}_\Phi(x).$$

So the global minimisers are exactly the assignments that maximise the number of satisfied clauses. We include this only as a short expressivity remark: unlike the earlier ordinary summed 2-adic MAX-CUT reduction [2], it uses a signed objective and assumes $p > 3$.

The corresponding positive encoding is unstable under added wells. Suppose the clause is instead encoded positively by the failure-sum

$$s_C(x) := \delta(\lambda_1; x) + \delta(\lambda_2; x) + \delta(\lambda_3; x)$$

and by the allowed set $S_C = \{0, 1, 2, 3\} \setminus \{3\}$. In the affine row $u_C^\top x \neq t_C$ above, this is the same missing-value phenomenon after moving the constants to the right-hand side. More generally, for $S \subseteq \{0, 1, 2, 3\}$, set

$$W_S(t) = \sum_{a \in S} |t - a|_p.$$

When $p > 3$ and $t \in \{0, 1, 2, 3\}$, this positive term is minimised exactly on S . Adding the missing positive well at the excluded value changes W_{S_C} into $W_{\{0,1,2,3\}}$, which is constant on Boolean inputs; doing this for every clause leaves only the unary Boolean pinning. The signed reward keeps the excluded affine value visible instead of erasing it by filling the missing well.

4 Sudoku as a special case

Sudoku has standard exact-cover, constraint-programming, and SAT formulations [3, 4, 5], and puzzle variants have their own complexity theory [6]. The construction below is not meant to compete with those solvers. Sudoku is used because its variables, domains, and all-different constraints are small enough to write out explicitly.

4.1 Sudoku

A (standard) Sudoku puzzle consists of a partially filled 9×9 grid. The task is to assign a digit in $\{1, \dots, 9\}$ to each empty cell such that:

1. each row contains all digits $1, \dots, 9$ exactly once;
2. each column contains all digits $1, \dots, 9$ exactly once; and
3. each 3×3 box contains all digits $1, \dots, 9$ exactly once.

Equivalently, in each unit (row, column, box) all entries are *pairwise distinct* and each entry lies in $\{1, \dots, 9\}$.

Standard Sudoku fits Corollary 3 directly. The vertices are the 81 cells, two vertices are adjacent when the corresponding cells share a row, column, or 3×3 box, unclued cells have domain $\{1, \dots, 9\}$, and clue cells have singleton domains. A proper list-colouring of this graph is exactly a valid Sudoku completion.

The Sudoku facts below are direct corollaries of the unary-well and edge-indicator lemmas. We state them before writing the concrete objectives.

4.2 Positive pinning: snapping coefficients to digits

Let $D = \{1, 2, \dots, 9\} \subset \mathbb{Z}_p$, and write the Sudoku variables as a vector

$$x \in \mathbb{Z}_p^{81},$$

with one coefficient per cell.

Definition 5 (Sudoku digit-snapping penalty). *For $p > 9$, define the digit-snapping penalty*

$$U_D(t) := \sum_{k=1}^9 |t - k|_p.$$

By Lemma 2 applied to the common domain D , one has $U_D(t) \geq 8$ for every $t \in \mathbb{Z}_p$, with equality if and only if $t \in D$.

This is the 9-valued analogue of the earlier “snap to $\{0, 1\}$ ” construction from [2]: the positive pinning observations $t = k$ for all $k \in D$ create a multiwell objective with minima exactly at D .

4.3 Negative signed residual terms: rewarding inequality via p-adic norms

The remaining Sudoku constraints are that entries in each row/column/box are pairwise unequal. Pairwise inequality can be encoded using differences $x_i - x_j$.

Corollary 4 (Sudoku differences as inequality indicators). *Let $p > 9$ and $a, b \in D$. Then*

$$|a - b|_p = \begin{cases} 0 & a = b, \\ 1 & a \neq b. \end{cases}$$

Proof. This is Lemma 1 with $q = 9$. □

4.4 Minimal and expanded residual objectives

Let $\mathcal{P}$ be the set of cell-pairs (i, j) that share a unit (same row, same column, or same box). Let $\mathcal{C}$ be the set of given clues; for $i \in \mathcal{C}$ let $g_i \in D$ be the fixed digit.

Specialising Definitions 3 and 4, together with Theorem 2, to the Sudoku graph gives the minimal synthetic dataset

$$L_{\min}(x) = \alpha \sum_{i \notin \mathcal{C}} \sum_{k=1}^9 |x_i - k|_p + \alpha \sum_{i \in \mathcal{C}} |x_i - g_i|_p - \sum_{(i,j) \in \mathcal{P}} |x_i - x_j|_p. \quad (1)$$

A separate coefficient on the pairwise term is unnecessary here: any positive edge weight can be scaled out, so we normalise that coefficient to 1. Thus α itself plays the role of the general pinning weight λ , and any choice $\alpha > 20$ satisfies the domination condition from Theorem 2. The minimal dataset therefore contains $9(81 - |\mathcal{C}|) + |\mathcal{C}| = 729 - 8|\mathcal{C}|$ positive observations and 810 negative observations. This is the formulation used in the correctness statements below.

For exposition, we also use the redundant expanded objective

$$L(x) = \alpha \sum_{i=1}^{81} \sum_{k=1}^9 |x_i - k|_p + \alpha_{\text{clue}} \sum_{i \in \mathcal{C}} |x_i - g_i|_p - \sum_{(i,j) \in \mathcal{P}} |x_i - x_j|_p. \quad (2)$$

Equation (2) keeps the nine digit-well residuals visible even on clue cells and then adds the clue pins separately. It therefore contains $729 + |\mathcal{C}|$ positive observations and the same 810 negative observations. The first two sums are positive pinning terms, and the final sum is a negative pairwise reward term. Equivalently, counting peer pairs per unit gives $27 \binom{9}{2} = 972$ negative terms, but 162 of these pairs lie simultaneously in a row/column *and* a box, leaving 810 distinct pairs overall.

This expanded form needs an explicit clue-weight assumption. On a clue cell, the first sum in (2) has the same value 8α for every digit $x_i \in D$, so clue fidelity is enforced only by $\alpha_{\text{clue}} |x_i - g_i|_p$. The same domination argument applies to (2) when $\alpha_{\text{clue}} > 20$; the simplest choice is $\alpha_{\text{clue}} = \alpha = 21$. Under that extra condition, (2) is a redundant expanded form of (1).

Once x is digit-valued, the pairwise part is simply

$$- \sum_{(i,j) \in \mathcal{P}} |x_i - x_j|_p = -|\mathcal{P}| + \sum_{(i,j) \in \mathcal{P}} \mathbf{1}_{x_i = x_j},$$

For the minimal objective on domain-respecting states, and for the expanded objective on digit-valued, clue-consistent states, the remaining terms are constant. In those regimes the loss differs from the *deduplicated* peer-conflict count only by an additive constant. The heuristics below do not optimise that deduplicated peer objective directly: on the row-permutation state space they use a duplicated unit-scope column/box surrogate that counts conflicts separately inside each column and each box. That surrogate is described in Section 5 and Appendix A.

Corollary 5 (Sudoku completion as a special case). *Let G_{Sud} be the graph on the 81 cells of a Sudoku puzzle, with an edge between two cells exactly when they share a row, column, or 3×3 box. For a puzzle with clue set $\mathcal{C}$ and clue digits g_i , define*

$$D_i := \begin{cases} \{g_i\}, & i \in \mathcal{C}, \\ \{1, \dots, 9\}, & i \notin \mathcal{C}. \end{cases}$$

Assume $\alpha > 20$. Then every global minimiser of the minimal objective (1) is a clue-respecting digit assignment that minimises the peer-conflict count. In particular, if the Sudoku instance is satisfiable, the global minimisers are exactly the valid completions of the puzzle.

Proof. Each row, column, and box of Sudoku forms a clique of size 9 in G_{Sud} . Thus a proper list-colouring assigns pairwise distinct digits in each unit. Since the available values are exactly $\{1, \dots, 9\}$, pairwise distinctness in a unit is equivalent to containing every digit exactly once. The singleton domains enforce the clues. Apply Theorem 2. $\square$

For Sudoku, the bound is explicit. In the deduplicated peer graph, every cell has degree 20, so Theorem 2 requires $\alpha > 20$ for the minimal objective (1); $\alpha = 21$ is enough. For the expanded objective (2), the same claim requires $\alpha_{\text{clue}} > 20$, because the all-digit wells are constant on clue cells. If one instead counts row, column, and box overlaps separately, each cell has $8+8+8 = 24$ incident negative terms, and the safe unary bound becomes > 24 . That threshold belongs to the duplicated unit-scope objective, not to the deduplicated peer objective (1).

Why ultrametric residuals help with negative weights. An ordinary term such as $-(x_i - x_j)^2$ over $\mathbb{R}$ rewards large differences, not just nonzero differences, and needs separate bounds. An ultrametric residual distinguishes equality from inequality. The discrete metric has this property by definition, and for $p > 9$ the p -adic norm has the same behaviour on digit-valued differences: $|x_i - x_j|_p \in \{0, 1\}$.

5 Illustrative computations

These computations are illustrative only: the point is to see the signed construction searched by simple local rules, not to propose a competitive Sudoku solver.

To keep local updates cheap, the implementation does not minimise (2) directly. Instead, on the digit-valued, clue-consistent row-permutation state space it minimises the duplicated unit-scope column/box conflict count

$$H_{\text{cb}}(x) := \sum_{c=1}^9 \text{conf}(C_c; x) + \sum_{b=1}^9 \text{conf}(B_b; x),$$

where C_c and B_b are the cells in column c and box b , and

$$\text{conf}(U; x) := \sum_{d=1}^9 \binom{\#\{u \in U : x_u = d\}}{2}.$$

Because each row is kept as a permutation of $\{1, \dots, 9\}$, the corresponding row contribution is identically zero and is omitted. This duplicated unit-scope loss is close to, but not identical with, the deduplicated peer-conflict objective coming from $\mathcal{P}$: a conflicting pair that lies simultaneously in a column and a box is counted twice. For the seed-0 initialisation shown later in Figure 4, the deduplicated peer-conflict count is 47 whereas $H_{\text{cb}} = 56$, because nine conflicting pairs lie in both a column and a box.

We implemented two simple local-search procedures for H_{cb} : a greedy row-swap search and a discrete Zubarev walk. Appendix A records the algorithms, a locality lemma for row swaps, and traces on a standard puzzle.

On a small sample of randomly carved puzzles, both heuristics solved all 36 instances tested. The row-swap search used fewer steps throughout the sample: its median step counts were 537, 2784, and 7040 at 36, 30, and 26 clues, versus 4533, 8751, and 9911 for the Zubarev walk. Appendix B gives the experimental setup, full table, and a representative loss curve.

6 Representational parsimony

A standard SAT or exact-cover lift introduces Boolean variables

$$y_{r,c,d} \in \{0, 1\},$$

with the intended meaning that cell (r, c) contains digit d . This produces $9 \times 9 \times 9 = 729$ indicators before any clues are applied. The p -adic construction instead uses an 81-cell digit space. This variable-level comparison is the main parsimony claim: one digit-valued parameter per semantic object rather than nine Boolean indicators per cell.

The minimal objective (1) is already a single scalar loss on that state space. The theorem-side expanded objective (2) contains 729 digit-snapping residuals, $|\mathcal{C}|$ clue residuals, and 810 pairwise residuals. A structured constraint-programming model with 81 cell variables and 27 *alldifferent* constraints is already semantically close to the puzzle, so the comparison here is not based on raw constraint count alone. The signed p -adic construction packages the all-different structure into one residual objective on the 81 cell variables. The heuristics in Appendix A then evaluate the duplicated unit-scope surrogate H_{cb} on the restricted row-permutation state space.

Appendix C records an alternative encoding that maps digits to powers of two so that each row, column, or box contributes a single sum residual rather than pairwise terms. In local-search experiments on 50 benchmark puzzles, this encoding produced no successful solves: the search repeatedly stopped in local minima.

7 Conclusion

Negative weights in Euclidean squared-loss regression are usually treated as anti-shrinkage terms in selected directions. Here signed residuals have a different role. Positive synthetic terms pin variables to a finite set; after that, negative p -adic terms are bounded rewards for inequality on domain-respecting states.

The paper does not propose a competitive Sudoku solver. Sudoku is a finite-domain case study in which the variables, domains, and all-different constraints are explicit. The main result is the encoding theorem: finite-domain all-different systems admit polynomial-time signed p -adic residual encodings whose global minimisers are exactly the domain-respecting minimum-conflict assignments, and in the satisfiable case exactly the satisfying assignments.

Mihara’s work on p -adic manifold learning from impartial games gives an adjacent comparison [8]. There, 3-heap Nim is a benchmark for estimating a p -adic zero locus associated with the Grundy function from finite samples. The Sudoku example here does not infer a hidden game-theoretic set. It starts from known finite-domain constraints and compiles them into signed residual terms. The connection is benchmark design and interpretability, not the all-different theorem.

The same signed-residual construction is not limited to pairwise inequality. After Boolean pinning, a 3-CNF clause can be written as one forbidden affine equality, so signed terms can reward clause satisfaction one clause at a time. This connects p -adic residual objectives with

finite-domain optimisation beyond all-different constraints, while keeping the claim separate from solver performance.

Algorithm 1 Row-swap local search

Require: Puzzle with clues $\mathcal{C}$, max steps T , restarts R

```
1: for  $r = 1$  to  $R$  do
2:   Initialise each row as a permutation of  $1, \dots, 9$  consistent with its clues
3:   for  $t = 1$  to  $T$  do
4:     if no column/box conflicts then
5:       return solution
6:     end if
7:     Choose a row involved in a conflict
8:     Find the best swap of two non-clue cells in that row (largest decrease in  $H_{cb}$ )
9:     Apply the swap (or a random swap if no improving swap exists)
10:  end for
11: end for
12: return best assignment found
```

A Heuristic details

A.1 Row-swap local search

The theorem-side objective (2) induces an NP-hard global optimisation problem, consistent with the earlier NP-hardness result for p -adic regression. The heuristics, however, work directly on the digit-valued, clue-consistent row-permutation state space and minimise the duplicated unit-scope loss H_{cb} from Section 5. Concretely, we initialise each row as a random permutation of $1, \dots, 9$ consistent with its clues, and then apply row-wise swaps that reduce H_{cb} . This is a block-local search on the integer state space.

Lemma 4 (Locality of row swaps). *Let x be a digit-valued, clue-consistent Sudoku state in which each row is a permutation of $1, \dots, 9$. Let ξ swap two non-clue entries in a single row r , at columns $c_1 \neq c_2$. Then the change*

$$\Delta H_{cb}(\xi) := H_{cb}(x \oplus \xi) - H_{cb}(x)$$

is determined entirely by the conflict-pair contributions from columns c_1 and c_2 , together with the one or two boxes containing (r, c_1) and (r, c_2) . No other term of the digit-restricted objective changes.

Proof. On the row-permutation state space, the heuristic objective is exactly H_{cb} , so only column and box contributions matter. A within-row swap preserves the row permutation, so the omitted row contribution remains zero. Every cell outside row r is unchanged, and within row r only the two swapped positions change value. Hence only the two corresponding columns can change their conflict counts. Likewise, only the boxes containing those two cells can change. All other columns and boxes see the same multiset of digits before and after the swap, so their contributions are unchanged. $\square$

A.2 Trace on a standard puzzle

To make the row-swap heuristic concrete, Figure 3 shows the standard example puzzle used in many Sudoku tutorials. With seed 0 (restart 0), the row-wise random initialisation used by the solver produces the filled grid in Figure 4; this state has 56 column/box conflict pairs. Table 1 shows the first few swaps from that initial state.

Step 2 illustrates the locality claim. After Step 1, row 6 contains

$$(7, 3, 1, 8, 2, 5, 4, 9, 6),$$

5	3			7				
6			1	9	5			
	9	8					6	
8				6				3
4			8		3			1
7				2				6
	6					2	8	
			4	1	9			5
				8			7	9

Figure 3: Standard example puzzle used for the heuristic traces.

5	3	8	4	7	2	1	9	6
6	2	8	1	9	5	3	4	7
3	9	8	1	4	7	2	6	5
8	2	9	5	6	4	7	1	3
4	5	7	8	9	3	2	6	1
7	3	1	8	2	5	4	9	6
3	6	1	7	9	5	2	8	4
2	6	3	4	1	9	7	8	5
3	2	4	6	8	5	1	7	9

Figure 4: Seed-0 row-wise random initialisation (restart 0), before any swaps.

and the chosen move swaps the entries in columns 3 and 6, exchanging 1 and 5. By Lemma 4, only column 3, column 6, the box on rows 4–6 and columns 1–3, and the box on rows 4–6 and columns 4–6 need to be inspected. Before the swap these four units contribute

$$4 + 6 + 1 + 2 = 13$$

conflict pairs; afterwards they contribute

$$3 + 3 + 2 + 1 = 9.$$

Thus $\Delta H_{\text{cb}} = -4$, exactly as reported in Table 1. For this instance the solver reaches a full solution after 228 swaps on the same restart.

A.3 A discrete Zubarev walk

The row-swap heuristic above is tailored to the strong-snapping Sudoku setting. It is also natural to test a more generic training dynamic from the p-adic regression literature. In [7], Zubarev proposes a random-walk-based optimisation scheme for loss functions built from p-adic norms, precisely because such losses are locally constant almost everywhere and hence poorly served by gradient methods.

Step	Row	Swap (columns)	Swap (digits)	Type	Conflicts (before)	(after)
1	5	7 $\leftrightarrow$ 8	2 $\leftrightarrow$ 6	best	56	53
2	6	3 $\leftrightarrow$ 6	1 $\leftrightarrow$ 5	best	53	49
3	2	2 $\leftrightarrow$ 3	2 $\leftrightarrow$ 8	best	49	45
4	4	7 $\leftrightarrow$ 6	7 $\leftrightarrow$ 4	random	45	47
5	7	1 $\leftrightarrow$ 5	3 $\leftrightarrow$ 9	best	47	40
6	4	6 $\leftrightarrow$ 7	7 $\leftrightarrow$ 4	best	40	38
7	2	2 $\leftrightarrow$ 9	8 $\leftrightarrow$ 7	best	38	36
8	1	3 $\leftrightarrow$ 4	8 $\leftrightarrow$ 4	best	36	35

Table 1: First eight swap steps for the standard example puzzle (seed 0). The table records both the swapped column indices and the digits occupying those columns before each move. Rows and columns are 1-indexed.

Concretely, the proposed training dynamics is a discrete-time random process

$$w_{t+1} = w_t + \xi_t(w_t, \beta_t), \quad (3)$$

where $\beta_t > 0$ is a “continuity” (inverse temperature) parameter and ξ_t is sampled from a distribution biased toward loss-decreasing moves. In Zubarev’s formulation the conditional law has density (with respect to the p -adic Haar measure)

$$P(\xi \mid w, \beta) \propto \exp(-\beta(L(w + \xi) - L(w))). \quad (4)$$

Under mild conditions (in particular, that the set of improving moves has nonzero measure), Zubarev shows there is a regime of β values for which the expected loss decreases, i.e. $L(w_t)$ is a strict supermartingale [7].

In the digit-restricted implementation used here, the state is already digit-valued, clue-consistent, and row-permuted. We therefore treat the Zubarev walk as a *finite-move* stochastic optimiser:

- The parameter vector w is the current Sudoku filling $x \in D^{81}$.
- The loss is the duplicated unit-scope column/box count $H_{\text{cb}}(x)$.
- A move ξ is a swap of two non-clue cells within a single row (preserving each row as a permutation of $1, \dots, 9$).

For a chosen row r and current state x , let $\Xi_r(x)$ be the set of all such swaps, and let $\Delta H_{\text{cb}}(\xi) := H_{\text{cb}}(x \oplus \xi) - H_{\text{cb}}(x)$ be the change in loss after applying swap ξ (where $x \oplus \xi$ denotes “apply the swap”). Replacing the Haar measure in (4) with the counting measure on the finite set $\Xi_r(x)$ gives the discrete sampling rule

$$\Pr(\xi \mid x, \beta_t) = \frac{\exp(-\beta_t \Delta H_{\text{cb}}(\xi))}{\sum_{\xi' \in \Xi_r(x)} \exp(-\beta_t \Delta H_{\text{cb}}(\xi'))}. \quad (5)$$

When $\beta_t \rightarrow 0$ this is uniform random swapping; when $\beta_t \rightarrow \infty$ it concentrates on the best (most loss-decreasing) swap, recovering a soft version of the greedy stepwise algorithm. In practice we use a simple annealing schedule: β_t increases linearly from β_0 to β_1 over the allotted steps, gradually shifting from exploration to exploitation.

A.4 Zubarev trace on the standard puzzle

Using the same standard example puzzle, the Zubarev walk (seed 0, restart 0, $\beta_0 = 0.5$, $\beta_1 = 6.0$ linear schedule) starts from the same row-consistent initialisation with 56 column/box conflict pairs. Unlike the greedy stepwise method, the walk sometimes accepts loss-increasing or loss-neutral swaps early on; Table 2 shows the first few moves. For this instance the Zubarev walk reaches a full solution after 10,030 swaps on the same restart.

Algorithm 2 Zubarev walk over row swaps (heuristic)

Require: Puzzle with clues $\mathcal{C}$, max steps T , restarts R , schedule $\{\beta_t\}_{t=1}^T$

- 1: **for** $r = 1$ to R **do**
- 2: Initialise each row as a permutation of $1, \dots, 9$ consistent with its clues
- 3: **for** $t = 1$ to T **do**
- 4: **if** no column/box conflicts **then**
- 5: **return** solution
- 6: **end if**
- 7: Choose a row involved in a conflict
- 8: Enumerate row swaps $\Xi_r(x)$ and compute $\Delta H_{\text{cb}}(\xi)$ for each $\xi \in \Xi_r(x)$
- 9: Sample $\xi \in \Xi_r(x)$ with probability $\propto \exp(-\beta_t \Delta H_{\text{cb}}(\xi))$
- 10: Apply swap ξ
- 11: **end for**
- 12: **end for**
- 13: **return** best assignment found

Step	Row	Swap (columns)	Conflicts (before)	(after)
1	5	$2 \leftrightarrow 8$	56	56
2	6	$7 \leftrightarrow 8$	56	56
3	2	$3 \leftrightarrow 7$	56	55
4	8	$2 \leftrightarrow 8$	55	51
5	4	$2 \leftrightarrow 3$	51	50
6	4	$6 \leftrightarrow 7$	50	51
7	2	$7 \leftrightarrow 8$	51	52
8	8	$3 \leftrightarrow 7$	52	49

Table 2: First eight Zubarev-walk swap steps for the standard example puzzle (seed 0). Rows and columns are 1-indexed.

B Illustrative experiments

We implemented both heuristics above in pure Python (see accompanying code). The prime was set to $p = 11$, but in the digit-restricted regime only the induced indicator behaviour matters: the search evaluates the integer conflict count H_{cb} directly and never computes a p -adic valuation numerically, so the p -adic norm routines in the code are illustrative of the construction rather than part of the solve path.

B.1 Experimental setup

To obtain a quick, reproducible testbed, we generated random puzzles by:

1. sampling a random solved grid; then
2. *carving* it down to a specified clue count by removing entries uniformly at random.

This procedure does *not* enforce uniqueness of the resulting puzzle. Each carved puzzle was then solved with up to $T = 60,000$ steps and $R = 15$ random restarts. All results reported below use a fixed RNG seed (123) for reproducibility.

The table can be regenerated from the repository root with the following commands:

```
python3 code/run_experiments.py --outdir outputs/paper_stepwise \
--seed 123 --n 6 --clues 36,30,26 --max-steps 60000 \
--restarts 15 --method stepwise
```

Method	Clues	Puzzles	Solved	Median steps	Median time (s)
Stepwise	36	6	6	537	0.270
Stepwise	30	6	6	2784	1.315
Stepwise	26	6	6	7040	3.990
Zubarev walk	36	6	6	4533	2.091
Zubarev walk	30	6	6	8751	3.951
Zubarev walk	26	6	6	9911	4.786

Table 3: Heuristic performance on randomly carved puzzles (not uniqueness-checked), using $T = 60,000$ steps and $R = 15$ restarts. The Zubarev walk uses a linear schedule $\beta_0 = 0.5$ to $\beta_1 = 6.0$.




Figure 5: Conflict trajectory for the stepwise row-swap heuristic on the fast-carved 30-clue instance with puzzle seed 955101029 (solve seed 2638744256), from the row-wise random initialisation on restart 0.

```
python3 code/run_experiments.py --outdir outputs/paper_zubarev \
--seed 123 --n 6 --clues 36,30,26 --max-steps 60000 \
--restarts 15 --method zubarev --beta0 0.5 --beta1 6.0 \
--beta-schedule linear
```

B.2 Results

Table 3 summarises 36 runs (six puzzles at each clue count, for each heuristic). Both solvers succeeded on all instances in this small sample; runtimes are machine-dependent and noisy because the puzzles are randomly carved, not difficulty-normalised, and the sample is too small to support stronger algorithmic claims.

Figure 5 shows the conflict trajectory for the stepwise row-swap heuristic on the specific fast-carved 30-clue instance with puzzle seed 955101029 (solve seed 2638744256), using the row-wise random initialiser described above, $T = 60,000$, and $R = 15$. The plotted run is restart 0, and the vertical axis is the duplicated column/box conflict count H_{cb} .

Experiment	Setting	Solved	Summary
E1	Greedy cell-swap on 15 primes, 50 puzzles, 3 random initialisations per puzzle–prime pair (2250 runs)	0/2250	Mean lift 0.003; mean 12.5 steps
E2 greedy	Greedy cell-swap on the 5 strongest E1 primes, again with 50 puzzles and 3 initialisations (750 runs)	0/750	Mean lift 0.009; mean 19.3 steps
E2 SA	Simulated annealing on the same 750 runs	0/750	Mean lift 0.000; mean 5000 steps

Table 4: Archived local-search results for the powers-of-two encoding. The E1 prime sweep used primes 2, 3, 5, 7, 11, 13, 17, 23, 47, 73, 97, 127, 257, 521, 2311.

C Powers-of-two experiments

An alternative maps digits $d \in \{1, \dots, 9\}$ to powers of two,

$$d \mapsto 2^{d-1} \in A := \{1, 2, 4, \dots, 256\}.$$

If every cell variable is restricted to A , then the condition that a row, column, or box contains each digit exactly once is equivalent to a single sum constraint

$$\sum_{(r,c) \in g} \beta_{r,c} = 1 + 2 + \dots + 256 = 2^9 - 1 = 511,$$

for each unit g . This replaces the pairwise penalties by 27 residuals of the form

$$r_g(\beta) := 511 - \sum_{(r,c) \in g} \beta_{r,c}.$$

For a prime p , one can then study the objective

$$L_p(\beta) := \sum_g |r_g(\beta)|_p,$$

optionally together with a regulariser that encourages $\beta_{r,c} \in A$ and enforces the clues.

Archived local-search experiments with this encoding were run on 50 Euler Project Problem 96 puzzles. Table 4 summarises the outcomes. None of these runs found a valid completion.

Funding

This research was supported by an Australian Government Research Training Program Scholarship.

Compliance with ethical standards

This article does not contain any studies with human participants or animals performed by the author.

Conflict of interest

The author declares that he has no conflict of interest.

Author contributions

This is a single-author paper. The author is responsible for the conceptual development, proofs, implementation, and writing.

Code and data availability

The Python implementation used for the heuristic experiments is publicly available at:

<https://github.com/solresol/sudoku-padic-regression>
Dataset archive: Hugging Face dataset
Revision: `ba49a8b2e51670f66db624d814d3f31678881b93`
Scripts: `code/padic_sudoku_regression.py`, `code/run_experiments.py`

The submitted source package includes the experiment CSV files, puzzle seeds, trace puzzle and solution files, and figure source used for Table 3 and Figure 5.

References

- [1] G. Baker, S. McCallum, and D. Pattinson. Linear regression in p -adic metric spaces. *p-Adic Numbers, Ultrametric Analysis and Applications* 17(4), 333–347 (2025). <https://doi.org/10.1134/S2070046625040016>
- [2] G. Baker. NP-hardness of p -adic linear regression. Accepted for publication in *p-Adic Numbers, Ultrametric Analysis and Applications* (2026); preprint at arXiv:2602.13278. <https://arxiv.org/abs/2602.13278>
- [3] Donald E. Knuth. *Dancing Links*. *Millennial Perspectives in Computer Science*, 187–214, 2000. Also available as arXiv:cs/0011047. <https://doi.org/10.48550/arXiv.cs/0011047>
- [4] Helmut Simonis. *Sudoku as a Constraint Problem*. In *Modelling and Reformulating Constraint Satisfaction Problems*, Fourth International Workshop, Sitges (Barcelona), Spain, October 1, 2005. <https://citeseerx.ist.psu.edu/document?doi=4f069d85116ab6b4c4e6dd5f4776ad7a6170faaf>
- [5] Inês Lynce and Joël Ouaknine. *Sudoku as a SAT Problem*. In *Proceedings of the Ninth International Symposium on Artificial Intelligence and Mathematics (ISAIM 2006)*, 2006. <https://people.mpi-sws.org/~joel/publications/sudoku05abs.html>

- [6] T. Yato and T. Seta. Complexity and completeness of finding another solution and its application to puzzles. *IEICE Transactions on Fundamentals*, 86-A(5):1052–1060, 2003.
- [7] A. P. Zubarev. p -Adic polynomial regression as alternative to neural network for approximating p -adic functions of many variables. *p-Adic Numbers, Ultrametric Analysis and Applications* 17(4), 413–420 (2025). <https://doi.org/10.1134/S2070046625040065>
- [8] T. Mihara. p -adic manifold learning and benchmark tasks from impartial games. arXiv:2605.04374 (2026). <https://arxiv.org/abs/2605.04374>
- [9] T. A. Hua and R. F. Gunst. Generalized ridge regression: a note on negative ridge parameters. *Communications in Statistics - Theory and Methods* 12(1), 37–45 (1983). <https://doi.org/10.1080/03610928308828440>
- [10] D. Kobak, J. Lomond, and B. Sanchez. The optimal ridge penalty for real-world high-dimensional data can be zero or negative due to the implicit ridge regularization. *Journal of Machine Learning Research* 21(169), 1–16 (2020). <https://www.jmlr.org/papers/v21/19-844.html>
- [11] P. Patil, J.-H. Du, and R. Tibshirani. Optimal ridge regularization for out-of-distribution prediction. In *Proceedings of the 41st International Conference on Machine Learning, Proceedings of Machine Learning Research* 235, 39908–39954 (2024). <https://proceedings.mlr.press/v235/patil24a.html>
- [12] A. Wiesel, Y. C. Eldar, and A. Yeredor. Linear regression with Gaussian model uncertainty: Algorithms and bounds. *IEEE Transactions on Signal Processing* 56(6), 2194–2205 (2008). <https://doi.org/10.1109/TSP.2007.914323>
- [13] A. J. Cannon. Negative ridge regression parameters for improving the covariance structure of multivariate linear downscaling models. *International Journal of Climatology* 29(5), 761–769 (2009). <https://doi.org/10.1002/joc.1737>